

Homepage: <https://chunhuizng.github.io>
Email: chunhui.zhang.gr@dartmouth.edu

GitHub: github.com/chunhuizng
Phone: 1-781-652-1380

Research Focus	RL Post-Training and Multimodal Alignment on Language, vision, and audio		
Education	Dartmouth College Ph.D. candidate in Computer Science Advisor: Professor Soroush Vosoughi	Hanover, NH, USA	Aug. 2023–Feb. 2026 (expected)
	Brandeis University M.S. in Computer Science GSAS Research Fellowship	Waltham, MA, USA	Sep. 2021–Jun. 2023
	Northeastern University B.S. in Computer Science Outstanding Honor Thesis Award	China	Sep. 2017–Jun. 2021
Experience	Amazon Applied Scientist Intern <i>RL Post-Training for multi-modal Computer-Use Agents (CUA)</i> Scaled vision–language–action models (7B) training for adaptive-thinking Graphical User Interface (GUI) agents, enabling adaptive reasoning through multi-stage RL via verl and vLLM. Interleaved on-policy RL with SFT on ground-truth hard trajectories to bridge OOD gaps and stabilize visual grounding in GUI environments.	Seattle, WA, USA	Sep. 2025 - Dec. 2025
	Google DeepMind Research Intern <i>Gemma 3n Post-Training for Real-time Streaming</i> Developed a synthetic data framework that eliminates reliance on human-crafted streaming labels while maintaining near-identical accuracy to non-streaming models. Built a full-stack RL+SFT framework for audio LLMs that achieves SOTA in real-time streaming ASR/AST. Engineered a high-throughput rollout-training pipeline in PyTorch and SGLang for Gemma 3n audio, accelerating RL (GRPO) by 5×. Contributed to Gemma 3n technical report and open-source support .	Mountain View, CA, USA	Jun. 2025–Sep. 2025

Selected Publications

Honda Research Institute USA	San Jose, CA, USA
Research Intern	Jun. 2024–Sep. 2024
<i>Multimodal LLM Post-Training (8B–70B) for Social Reasoning</i>	
Published a <i>spotlight</i> paper at ICML 2025 (Top 2.59%). Filed a US patent. Developed long-context multimodal LLMs and accelerated inference-time scaling using vLLM framework. Enabled long-context modeling of human behavior and autonomous driving in multimodal environments; scalable to LLMs up to 70B/405B parameters.	

Overcoming Multi-step Complexity in Theory-of-Mind Reasoning: A Scalable Bayesian Planner

Chunhui Zhang, Zhongyu Ouyang, Sean Dae Houlihan, Kwonjoon Lee, Nakul Agarwal, Soroush Vosoughi, Shao-Yuan Lo
ICML 2025 (Spotlight, Top 2.59%)

[Paper](#) | [Code](#)

First scalable solution for multi-step Theory-of-Mind reasoning. Uses Bayesian inverse planning for global planning, then lets LLMs focus on local reasoning. Works on 70B+ models where others fail.

Growing Through Experience: Scaling Episodic Grounding in Language Models

Chunhui Zhang, Elsie Wang, Zhongyu Ouyang, Xiangchi Yuan, Soroush Vosoughi

ACL 2025 Main Conference (Oral Presentation, Top 3.24%)

[Paper](#) | [Code](#)

Post-trained reasoning LLMs (across 3B, 8B, 70B) on MCTS-sampled data from physical simulators.

Pretrained Image-Text Models are Secretly Video Captioners

Chunhui Zhang*, Yiren Jian*, Zhongyu Ouyang, Soroush Vosoughi
NAACL 2025 Main Conference (Oral Presentation, Top 2.88%)

[Paper](#) | [Code](#)

RL post-training recipe that achieved **Top-2** on PapersWithCode video captioning leaderboard, outperforming industry MLLM captioners.

Knowing More, Acting Better: Hierarchical Representation for Embodied Decision-Making

Chunhui Zhang, Zhongyu Ouyang, Zheyuan Liu, Soroush Vosoughi
EMNLP 2025 Findings

Refines vision-language-action LLM representations to enable more effective PPO-based RL training in embodied AI.

Working Memory Identifies Reasoning Limits in Language Models

Chunhui Zhang, Yiren Jian, Zhongyu Ouyang, Soroush Vosoughi

EMNLP 2024 Main Conference | Paper | Code

Introduced working memory as a diagnostic tool for LLM reasoning limits. This work inspired a follow-up NAACL paper on long-context multimodal understanding.

Temporal Working Memory: Query-Guided Segment Refinement for Enhanced Multimodal Understanding

Xingjian Diao*, **Chunhui Zhang***, Weiyi Wu, Zhongyu Ouyang, Peijun Qing, Ming Cheng, Soroush Vosoughi, Jiang Gui

NAACL 2025 Findings | Paper | Code

Inspired by working memory in my EMNLP 2024 paper, this work is a follow-up study on long-context video-language understanding.

Under Review

Model Priors Shape Experience: RL for Complex Audio Long-form Reasoners

Chunhui Zhang

[Code](#)

First to implement Qwen2.5-Omni and vLLM for faster RL reasoning across audio and other unified modalities.

Selected Collaborations

Superficial Self-Improved Reasoners Benefit from Model Merging

Xiangchi Yuan, **Chunhui Zhang**, Zheyuan Liu, Dachuan Shi, Soroush Vosoughi, Wenke Lee

EMNLP 2025 | Paper | Code

SoundMind: RL-Incentivized Logic Reasoning for Audio-Language Models

Xingjian Diao, **Chunhui Zhang**, Keyi Kong, Weiyi Wu, Chiyu Ma, Zhongyu Ouyang, Peijun Qing, Soroush Vosoughi, Jiang Gui

EMNLP 2025 (Oral Presentation)

Is It Navajo? Accurate Language Detection in Endangered Athabaskan Languages

Ivory Yang, Weicheng Ma, **Chunhui Zhang**, Soroush Vosoughi

NAACL 2025 (Oral Presentation, Top 2.88%) | Paper | Code

Learning Musical Representations for Music Performance Question Answering

Xingjian Diao, **Chunhui Zhang**, Tingxuan Wu, Ming Cheng, Zhongyu Ouyang, Weiyi Wu, Jiang Gui
EMNLP findings 2024

Expedited Training of Visual Conditioned Language Generation via Redundancy Reduction

Yiren Jian, Tingkai Liu, Yunzhe Tao, **Chunhui Zhang**, Soroush Vosoughi, Hongxia Yang
ACL 2024 (Oral Presentation, Top 3.10%)

Aligning Relational Learning with Lipschitz Fairness

Yanling Jia*, **Chunhui Zhang***, Soroush Vosoughi
ICLR 2024 – Note: Co-first author Jia was a master student I mentored.

Mitigating Emergent Robustness Degradation on Graphs while Scaling-up

Xiangchi Yuan*, **Chunhui Zhang***, Yijun Tian, Yanfang Ye, et al.
ICLR 2024 – Note: Co-first author Yuan was a master student I mentored.

When Sparsity Meets Contrastive Models: Less Data Can Bring Better Class-Balanced Representations

Chunhui Zhang, Chao Huang, Yijun Tian, Qianlong Wen, et al.
ICML 2023 – AAAI-DCAA 2023 Best Paper Runner-up Award

Chasing All-Round Graph Representation Robustness: Model, Training, and Optimization

Chunhui Zhang, Yijun Tian, Mingxuan Ju, Zheyuan Liu, et al.
ICLR 2023

Look Twice as Much as You Say: Scene Graph Contrastive Learning for Self-Supervised Image Caption Generation

Chunhui Zhang, Chao Huang, Youhuan Li, Xiangliang Zhang, Yanfang Ye, et al.
CIKM, 2022.

Full publication list including 30+ papers at ICML, NeurIPS, ICLR, ACL, NAACL, EMNLP, WWW (TheWebConf), KDD, CIKM, ICDM, ICWSM, etc. available at chunhuizng.github.io

Honors and Awards	ICML Spotlight Recognition (Top 2.59%)	2025
	ACL Oral Presentation Award (Top 3.24%)	2024, 2025
	NAACL Oral Presentation Award (Top 2.88%)	2025
	AAAI-DCAA Best Paper Runner-up Award	2023
	Graduate School of Arts and Sciences Fellowship	2021–2023
Service	Program Committee/Reviewer: ICML, ICLR, NeurIPS, ACL, NAACL, EMNLP, UAI, AISTATS, AAAI (2023–2025) Journal Reviewer: TMLR, IEEE TKDE, IEEE TNSE, ACM TIST, Neurocomputing	